

The fmlogit Package: A Document

Xinde James Ji

Thursday, May 19, 2016

This document provides documentations for the fmlogit package in R. Updates will be published at my github site: <https://github.com/flkidd/fmlogit>. Any suggestions or concerns are welcomed¹.

Motivation

Fractional multinomial responses arises naturally in various occasions. For example, a municipality allocates its budgets to multiple departments, and we are interested in the proportion of the budgets that each department receives. Or, there are multiple candidates in a presidential election, and we are interested in the percentage of support for each candidate in each state.

However, fractional multinomial logit model is *underrepresented* in the booming era of statistical softwares. The model itself is a coupling of *fractional response models*, which deals with responses which are proportional or fractional, and *multinomial response models*, which deals with binary responses of multiple options. As there are multiple softwares that deals with fractional logits and multinomial logits, the only software package that deals with fractional multinomial logits is Stata's fmlogit package by Maarten Buis. This package will contribute to the sea of software packages by providing an implementation of the model in R.

Note here that fractional multinomial logit model is a consistent estimator of fractional multinomial responses. Other estimation strategies are certainly useful in estimating fractional multinomials. Dirichlet regression models and two-limit tobit can also deal with responses that are ranged between 0 and 1. However, those models do not have the additional restriction that the proportions sums up to one, and thus will not be preferable in this case.

Econometric Model

The basis of this function is Papke and Wooldridge(1996)'s paper, in which they proposed a quasi-maximum likelihood(QMLE) estimator for fractional response variables. As their approach applies to binary response variables, here we expand it to a multinomial response variables with fractional structure.

We start by writing:²

$$E(y_{ij}|x_i) = G(x_i\beta_j)$$

for the j^{th} choice of the i^{th} observation, where $G(\cdot)$ is a know function satisfying $0 < G(z) < 1$ for all $z \in \mathbb{R}$. Note that here we only allow for common covariates of x_i , and not for choice-specific attributes. Following the logit convention, $G(\cdot)$ is chosen to be the multinomial logit function, with the form:

$$G(z_j) = \frac{\exp(z_j)}{\sum_{k=1}^J \exp(z_k)}$$

And the multinomial likelihood function, is thus given by

$$\ln(L_i) = \sum_{j=1}^J y_{ij} \ln(G(x_i\beta_j))$$

¹email: xji@vt.edu

²The demonstration below is in individual specific notation, but matrix notation is not hard to obtain from the individual specific notations. The actual function uses matrix calculation, which increases algorithm speed.

And Papke and Wooldridge(1996) showed that the QMLE estimator of β , obtained by the the maximization problem

$$\operatorname{argmax}_{\beta} \sum_{i=1}^N \ln(L_i)$$

is a consistent estimator for β if $G(z)$ is the correct functional form for $E(y|x)$.

To estimate the standard error for the QMLE estimator, define $g(z_j) \equiv \partial G(z_j)/\partial z_j$, the partial derivative of the multinomial logit function with respect to choice j . Specifically, $g(z_j)$ has the following functional form:

$$g(z_j) = \frac{\hat{E}\hat{S} - \hat{E}^2}{\hat{S}^2}$$

where $\hat{E} = \exp(x_i\beta_j)$, and $\hat{S} = \sum_{k=1}^J \exp(x_i\beta_k)$.

A robust asymptotic standard error , $\hat{\beta}_j$, is given by the diagonal element of the following matrix:

$$\hat{A}_j^{-1} \hat{B}_j \hat{A}_j^{-1}$$

where

$$\hat{A} = \sum_{i=1}^N \frac{\hat{g}_{ij}^2 \mathbf{x}_i' \mathbf{x}_i}{\hat{G}_{ij}(1 - \hat{G}_{ij})}$$

$$\hat{B} = \sum_{i=1}^N \frac{\hat{u}_{ij}^2 \hat{g}_{ij}^2 \mathbf{x}_i' \mathbf{x}_i}{[\hat{G}_{ij}](1 - \hat{G}_{ij})^2}$$

in which \hat{u}_{ij} is the residual for the j^{th} choice of the i^{th} observation, given by $\hat{u}_{ij} = y_{ij} - G(x_i\beta_j)$. Specifically, \hat{A} is the information matrix, which is not a consistent estimator itself, and \hat{B} is a weight correction for A .

In most binary / multinomial response models, the convention is to treat one of the choices as a baseline. Here we apply the same logic, and treat $j=1$ as the baseline scenario. This implicitly generates a restriction that $\beta_1 = 0$, and all other betas are the marginal difference to the baseline case.

Marginal and Discrete Effects

Interpreting marginal and discrete effects for limited dependent variables can be tricky, and this is especially true for multinomial logit models. The coefficients obtained in the regression model represents the logit-transformed odds ratio for that specific choice against the baseline choice. This is not intuitive at all in terms of what are the actual effects on that specific choice. The bottom line is, the coefficients and standard errors obtained in the original models are not the basis for evaluating hypotheses.

Marginal Effects

Instead, researchers need to compute what's called the "marginal effect", as we usually do in linear models. The marginal effect of the multinomial models actually has a very distinctive form:

$$ME_{jk} = \frac{\partial p_j}{\partial x_k} = p_j(\beta_{kj} - \bar{\beta}_i)$$

where p_j is an $1 \times N$ vector of predicted probabilities for choice j , and $\bar{\beta}_i = \sum_{m=1}^J \beta_{km} p_m$ is the probability weighted average of β_{km} . This shows that the marginal effects among different individuals are actually different given different predicted probabilities of choice j .

Typically, two types of summary measures are used to illustrate the global average marginal effects. The first one is called marginal effects at the mean (MEM) in the code. In algebraic form, this is represented as

$$MEM_{jk} = \bar{p}_j(\beta_{kj} - \bar{\beta}_i)$$

where \bar{p}_j is the predicted value of choice j at the mean of all X covariates. Centering observations around the mean simplifies the calculation, however it ignores the potential heterogeneity in marginal effects, especially at the extreme values.

Another measure is called average marginal effects (AME). This can be written as:

$$AME_{jk} = \frac{1}{N} \sum_{i=1}^N p_j(\beta_{kj} - \bar{\beta}_i)$$

However, according to Greene(2003), there is no agreement as to which one is preferred. A more inclusive approach will be to plot the marginal effect of interest across all individuals. This is not provided in the function, but can certainly be implemented in a straightforward way in R.

Discrete Effects

Discrete effect is a little bit different from marginal effects. Instead of calculating the slope of the coefficients, discrete effect considers the impact of a discrete change in one covariates on the predicted outcome variables. This is especially useful for dummy variables, where calculating marginal effects does not make much sense.

The discrete effect has a straight-forward form. Consider a discrete change of a dummy variable k from 0 to 1. This is just

$$D_j = Pr(y = j | \mathbf{x}_{x_k=1}) - Pr(y = j | \mathbf{x}_{x_k=0})$$

The change in predicted value by setting $x_k = 1$ and $x_k = 0$.

Similar to the marginal effect case, we can calculate discrete effects at the mean (DEM) by predicting the outcome when all other covariates at the mean, or average discrete effects, which averages the predicted difference across all observations.

Standard Errors

Here we adopt the Krinsky-Robb method to compute standard errors for marginal and discrete effects. As oppose to the delta method commonly used in other programs such as Stata, Krinsky-Robb is a simulation-based method. The idea of Krinsky-Robb is that, to calculate the variance of a function $Var(f(\beta))$, we do the following step:

For i in $1:N$, where N is a very large number,

- 1) Sample from the known distribution of β
- 2) For each of the sample, calculate $f(\beta)$
- 3) Take the empirical variance of $f(\beta)$.

And after sufficiently large sample size, the empirical variance converges to the theoretical variance.

Note that the Krinsky-Robb sampling is done by setting all other covariates to the mean, i.e., the MEM and the DEM approach.

Practical Concerns

Optimization Method

This function calls *maxLik()* in package *maxLik* to maximize the quasi-likelihood function. The *maxLik* function is a wrapper which provides several different maximization methods, including most *optim()* methods in the base package, as well as other useful methods such as BHHH(Berndt-Hall-Hall-Hausman). The choice of optimization method can create vastly different parameter estimates. Here it is recommended that either conjugate gradients(CG), or Berndt-Hall-Hall-Hausman(BHHH) to ensure convergence. In limited testing scenarios, BHHH typically has the best performance in terms of convergence for large datasets, while CG is faster in computation speed for smaller, easy to converge datasets.

Robust Standard Error

It is worth noting that the robust standard error created in this function is consistently lower than that created in Stata's *fmlogit* package, typically by about 20%. However, the robust SE here is a consistent estimator following Papke and Wooldridge(1996)'s $\hat{A}_j^{-1} \hat{B}_j \hat{A}_j^{-1}$ estimator, so it is recommended that the number should be used with caution.

References

- Papke, L. E. and Wooldridge, J. M. (1996), Econometric methods for fractional response variables with an application to 401(k) plan participation rates. *J. Appl. Econ.*, 11: 619-632.
- Wulff, Jesper N. "Interpreting Results From the Multinomial Logit Model Demonstrated by Foreign Market Entry." *Organizational Research Methods* (2014): 1094428114560024.
- Greene, W. H. (2003). *Econometric analysis* (5th ed.). Upper Saddle River, NJ.: Prentice Hall